

Project 2: Simple Window-based Reliable Data Transfer

1 Goal

The purpose of this project is to use UDP Socket and C/C++ programming language to implement a reliable data transfer protocol.

2 Instructions

1. In this project, you will be implementing a simple congestion window-based protocol built on top of Selective Repeat protocol (not stop & wait, not stop & forward, and not go-back-N protocol) described in the textbook. You must write one program implementing the server side, and another program implementing the client side. Only C/C++ are allowed to use in this project.
2. The two programs will communicate using the User Datagram Protocol (UDP) socket, which does not guarantee data delivery.
3. The client program will also act as a client application, requesting a file from the server.
4. The client program will take the hostname and port number of the server, and the name of a file it wants to retrieve from the server as a command line arguments. For example:
to run the server type: *shell > server < portnumber >*
to run the client type: *shell > client < server_hostname > < server_portnumber > < filename >*
5. The client will first send a message to the server which includes the name of the file requested. If the file exists, the server will divide the entire file into multiple packets, and then add some header information to each packet before sending them to the client. If the file does not exist at server, the client should display 404 not found error.
6. Note that your programs will act as both a network application (file transfer program) as well as a reliable transport layer protocol built over the unreliable UDP transport layer.

3 Error Handling

Although using UDP does not ensure reliability of data transfer, the actual rate of packet loss or corruption in LAN may be too low to test your program. Therefore, we are going to emulate packet loss by using 'tc' command in linux.

- The file transmission should be completed successfully, unless the packet loss rate is 100
- The timer on both the client and the server should work correctly to retransmission the lost packet.
- You can use the following commands to adjust parameters of the emulation:
- To check the current parameters for the private network ('eth1'):


```
tc qdisc show dev eth1
```
- To change current parameters to loss rate of 20% and delay 100ms:


```
tc qdisc change dev eth1 root netem loss 20% delay 100ms
```

- To delete the network emulation:

```
tc qdisc del dev eth1 root
```

- If network emulation hasn't yet setup or you have deleted it, you can add it to, e.g., 10% loss without delay emulation:

```
tc qdisc add dev eth1 root netem loss 10%
```

- You can also change network emulation to re-order packets. The command below makes 4 out of every 5 packets (1-4, 6-9, ...) to be delayed by 100ms, while every 5th packet (, 10, 15, ...) will be sent immediately:

```
tc qdisc change dev eth1 root netem gap 5 delay 100ms
```

- or if you're just adding the rule

```
tc qdisc add dev eth1 root netem gap 5 delay 100ms
```

- More examples can be found in [Network Emulation tutorial]
(<http://www.linuxfoundation.org/collaborate/workgroups/networking/netem>)

4 Hints

The best way to approach this project is in incremental steps. Do not try to implement all of the functionality at once.

- First, assume there is no packet loss. Just have the server send a packet, the client respond with an ACK, and so on.
- Second, introduce a large file transmission. This means you must divide the file into multiple packets and transmit the packets based on the current window size.
- Third, introduce packet loss. Now you have to add a timer on each sent and unAcked packet. If a timer times out, the corresponding (lost) packet should be retransmitted for the successful file transmission.

The credit of your project is distributed among the required functions. If you only finish part of the requirements, we still give you partial credit. So please do the project incrementally.

5 Requirements

- You must use an UDP socket for both server and client.
- You should print messages to the screen when the server or the client is sending or receiving packets. There are four types of output messages and should follow the formats below.

– Server: Sending packets

“Sending packet” [Sequence number] [WND] (“Retransmission”) (“SYN”) (“FIN”)

Example:

```
Sending packet 5095 5120 SYN
```

```
Sending packet 5096 5120
```

```
Sending packet 6120 5120
```

```
Sending packet 7144 5120
```

```
Sending packet 5096 5120 Retransmission
```

```
Sending packet 8168 5120 FIN
```

- Server: Receiving packets
“Receiving packet” [ACK number]

Example:

Receiving packet 5096
Receiving packet 6120

- Client: Sending packets
“Sending packet” [ACK number] (“Retransmission”) (“SYN”) (“FIN”)

Example:

Sending packet SYN
Sending packet 5096
Sending packet 6120
Sending packet 7144
Sending packet 6120 Retransmission
Sending packet 8168 FIN

- Client: Receiving packets
“Receiving packet” [Sequence number]

Example:

Receiving packet 5096
Receiving packet 6020

Such messages will be helpful for you to debug the programs, and we can use them to examine the correctness of your programs.

- The maximum packet size is 1024 bytes including a header.
- The window size must be in **the unit of bytes**. (Not in the unit of packet count) For example, if the window size is 5120 bytes, then five packets can be transmitted simultaneously.
- The sequence number is given in the unit of bytes as well. The maximum sequence number should correspond to 30 Kbytes (30720 bytes). You have to reset back the sequence number when it reaches the maximum value.
- Packet retransmission should be triggered when the timer times out.
- Here are the default values for some variables.
 - Maximum packet length (including all your headers): 1024 bytes
 - Maximum sequence number: 30720 bytes
 - Window size: 5120 byte
 - Retransmission time out value: 500 ms
- The client should save file to ‘received.data’ in the current working directory.
- After server finishes transmission, it should terminate the connection using FIN/FIN-ACK procedure. The client should implement TIME-WAIT mechanism, e.g., using 2*RTO as a waiting time.

6 Congestion Control: Extra Credit

If you implement TCP congestion control features then you can get up to 5% extra credit. Here are requirements for the congestion control.

- Create separate execution files for both server and client to show congestion control.
- Set initial slow start threshold as 15360 bytes.
- Set initial (congestion) window size as 1024 bytes.
- You should print slow start threshold when the server sends data packets.
(e.g. *Sending packet 5095 5120 **15360***)
- Following congestion control features will be checked for grading
 - Slow start
 - Handle packet loss (CWND and SSThresh)
 - Fast retransmission and fast recovery
 - Congestion avoidance

7 Due Date and Demo

1. You are required to demo your program on 3/19 and 3/20.
2. Submit an electronic copy of the project on CCLE on **11:50pm** 3/16.
3. Late submission is not allowed!

7.1 Demo

In the demo, we provide you the provided VM machine. You then use make to compile your programs, run your programs to deliver a test file from the server side to the client side. It is required by your program to print out the operations, which explain the delivery process on the screen. We may ask you to use different values for 'tc' command to test your programs. We will ask you to compare the received test files with the original one using the tool "diff".

Demo Procedure:

1. Sign up for Demo : TA will distribute the demo signup sheet in the discussion section of the 8th week.
2. In Demo : You need to prepare 3 slides to present. 1 slide for design and implementation, 1 slide for experiences you gain, 1 slide for lesson learn from project and suggestion to project.
 - TA will ask you to demo the function step by step
 - TA will also ask you questions during demo, you need to answer the question clearly. All question will related to your project implementation.
3. Demo time is 15 minutes.

8 Project Submission

1. Put all your files into a directory, must called *project2_UID1_UID2.tar*. UID is the student id of one of the students of the group.
2. Submit the file *project2_UID1_UID2.tar* on CCLE course website.
3. The *project2_UID1_UID2.tar* should contain the following files
 - Source codes (can be multiple files)
 - A report file (.doc or .pdf) no more than 3 pages. The report will contain:
 - Student names and Student IDs at the very beginning (**2 students per-group**).
 - Implementation description (header format, messages, timeouts, window-based protocol etc).
 - Difficulties that you faced and how you solved them.
 - Makefile
 - README file that includes your group information. The README will contain:
 - Group number, students name, and UID.
 - Brief description about how you divide the workload.
4. The TAs will only type “make” to compile your code, make sure your Makefile works in the provided VM machine.
5. Each group just needs to submit one set of files.